

LongListBench: A Benchmark for Long-List Entity Extraction Under Layout and OCR Noise

Anton Fedoruk^{1*} Serhii Shchoholiev^{1†} Akhil Mehta^{1‡}

¹Kay.ai, Brooklyn, NY, USA

January 22, 2026

Abstract

Existing document extraction benchmarks focus on key-value extraction, leaving long-list entity extraction underexplored. Yet business documents such as loss runs, invoices, and itemized bills commonly contain dozens to hundreds of repeated records. We introduce LongListBench, a synthetic benchmark for long-list extraction that pairs ground truth JSON with rendered PDFs and OCR transcripts. The benchmark injects seven document phenomena observed in production: page breaks, multi-row entities, duplicates, large documents, irrelevant tables, multi-column layouts, and merged cells. Across 80 documents (6,828 incident rows), zero-shot LLM baselines achieve 81.9% (Gemini 2.5), 80.0% (GPT-4o), and 78.1% (GPT-5.2) field-level F1, with the table format and layout disruptions posing the greatest challenges.

1 Introduction

Long-list entity extraction—recovering dozens to hundreds of repeated records from semi-structured documents—is a core requirement for document automation in domains such as insurance, finance, and procurement. While recent advances in document understanding models (e.g., layout-aware pretraining [1] and OCR-free approaches [2]) and general-purpose LLMs have improved extraction quality, robust evaluation of long-list scenarios remains limited.

Many established benchmarks focus on key-value style extraction or relatively short, form-like documents (e.g., FUNSD [3]) or narrow document types such as receipts (SROIE [4]). More recent datasets such as DocILE [5] include business documents and line items, but long lists in the wild often exhibit additional failure modes: repeated entities, page breaks, multi-column reading order, irrelevant tables, and table constructs such as merged cells.

*anton@kay.ai

†serhii@kay.ai

‡akhil@kay.ai

VRDU [6] highlights that hierarchical and long-list fields remain challenging for LLM-based extraction.

We introduce LongListBench, a benchmark designed to stress-test long-list extraction from semi-structured business documents with long incident/line-item lists under systematically injected document phenomena and OCR noise. The benchmark is inspired by recurring patterns observed in real-world claims documents.

1.1 Background and Motivation

This work originates from production challenges encountered at Kay.ai and in prior industry experience. A client engagement required generating Statements of Values (SOVs) from loss-run PDFs containing insurance claims—documents spanning hundreds of pages with varied formats and numerous layout artifacts. After evaluating off-the-shelf models and commercial extraction services, we identified long-list extraction as an underserved problem: existing tools performed adequately on short forms but degraded on documents with dozens to hundreds of repeated records. Similar challenges arose when processing itemized medical bills containing thousands of claim lines, exhibiting wide variation in table structures and OCR quality. These experiences motivated the development of a dedicated extraction pipeline (to be described in a subsequent paper) and, in turn, the need for a rigorous benchmark to measure progress on extraction methods that remain reliable as list length grows and as layout artifacts accumulate.

1.2 Research Questions

This work is organized around three practical questions:

- How do common long-list document phenomena (page breaks, duplicates, multi-row cells, multi-column layouts, irrelevant tables, merged cells) affect extraction quality?
- To what extent are end-to-end failures attributable to OCR transcription versus downstream extraction?
- How do strong off-the-shelf LLMs perform under a simple, reproducible zero-shot protocol?

1.3 Contributions

We make the following contributions:

- A reproducible benchmark generation pipeline that produces paired ground truth JSON, rendered PDFs, and OCR transcripts.
- A dataset of 80 documents (40 detailed, 40 table) containing 6,828 incident rows across four difficulty tiers, with an extreme tier reaching 500 incidents per document.
- A taxonomy of seven injected problem types and evaluation scripts for field-level scoring and OCR identifier coverage.

- Baseline results for GPT-5.2, GPT-4o, and Gemini 2.5 under a shared prompt, highlighting remaining gaps in long-list extraction.

2 Related Work

Research on information extraction (IE) from visually rich documents has produced a broad ecosystem of datasets and models. However, much of the public evaluation landscape emphasizes either short documents (e.g., forms) or key-value extraction, leaving long-list entity extraction underexplored.

2.1 Document IE benchmarks

Early and widely used benchmarks such as FUNSD [3] focus on form understanding in noisy scans. Receipt datasets and challenges such as SROIE [4] emphasize OCR and key fields in narrow document types. These benchmarks are valuable, but typically contain relatively short documents and do not directly stress long lists of repeated entities.

DocILE [5] broadens the scope to business documents and includes line-item recognition, which is closer in spirit to long-list extraction. VRDU [6] further argues that hierarchical and repeated fields (e.g., invoice line items) remain difficult for LLM-based extraction. Our benchmark complements these efforts by focusing on list length, repeated entity boundaries, and a targeted taxonomy of long-list failure modes.

2.2 Document understanding models

Layout-aware pretraining approaches such as LayoutLM [1] jointly model textual content and 2D document structure, yielding strong performance on a range of document understanding tasks. In parallel, OCR-free approaches such as Donut [2] avoid explicit OCR by directly generating structured outputs from document images, mitigating OCR error propagation at the cost of specialized training.

In contrast, our work is model-agnostic: we provide paired PDF, OCR transcript, and ground truth, enabling evaluation of OCR-based pipelines, OCR-free models, and LLM-based extraction. Our primary goal is to support reproducible measurement of long-list extraction robustness under realistic layout artifacts.

3 Benchmark Construction

We construct LongListBench, a synthetic benchmark for long-list entity extraction in semi-structured business documents with long incident/line-item lists, inspired by recurring patterns observed in real-world claims documents. Each benchmark instance consists of (i) structured ground truth incidents (JSON), (ii) a rendered PDF, and (iii) an OCR transcript of the PDF in Markdown.

3.1 Entity schema

Ground truth incidents follow the schema in `benchmarks/models/loss_run.py`, represented as a Pydantic model. The schema includes incident identifiers, policy metadata, narrative text, and nested financial breakdowns (`bi`, `pd`, `lae`, `ded`). While downstream workflows often emphasize fields such as `incident_number`, `company_name`, `date_of_loss`, `status`, `driver_name`, `coverage_type`, and `total_incurred`, our evaluator requires and scores the full schema.

3.2 Document generation

Benchmark instances are generated with seeded randomness for reproducibility (`benchmarks/generate_claims_benchmark.py`). Structured incidents are created by `benchmarks/synthetic/generate_claim_data.py`. We then render the incidents into visually rich documents using `benchmarks/synthetic/generate_html.py` and one of two layouts:

1. **Detailed**: a repeated incident block with narrative text and a small financial table.
2. **Table**: a compact tabular representation formatted as a CSV-like table.

The HTML is rendered into PDF using headless Chromium via Playwright (`benchmarks/synthetic/html_to_pdf.py`).

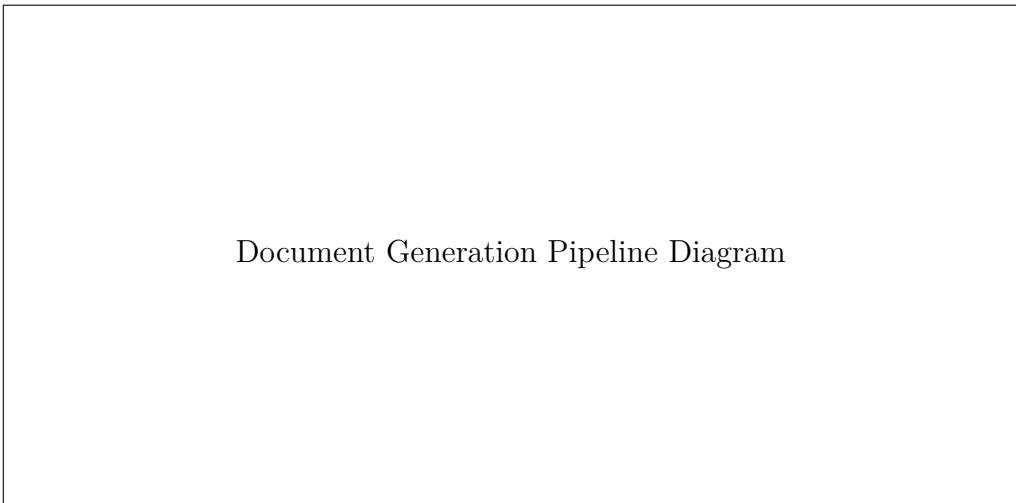


Figure 1: Overview of the benchmark document generation pipeline.

3.3 Injected problem types

We inject seven recurring document phenomena that complicate long-list extraction. These effects are applied at the HTML level prior to PDF rendering.

Page breaks. In real-world documents, page boundaries frequently split logical entities mid-record. A single incident may begin on one page with identifiers and description, while its financial breakdown appears on the next. This is particularly common in loss runs and itemized bills where dense formatting leaves no natural breakpoints. OCR systems typically process pages independently, producing separate text blocks that must be reassembled. Extraction models must recognize continuation patterns and avoid treating the second half of a split record as a new entity or dropping it entirely.

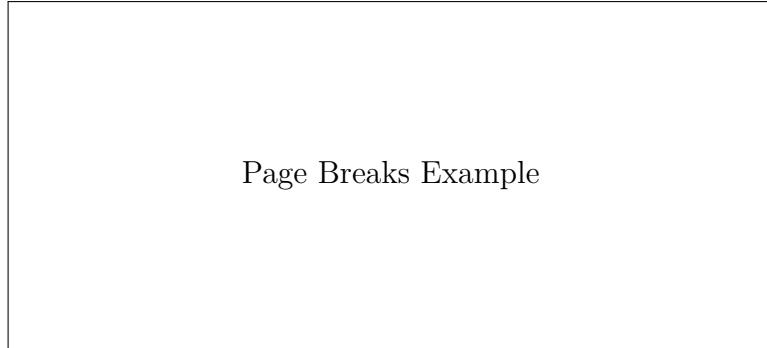


Figure 2: Example of an incident split across a page boundary.

Multi-row entities. Table cells often contain text that wraps across multiple lines, especially for description fields, addresses, or claimant lists. When OCR linearizes such content, it may interleave text from adjacent columns or treat each line as a separate cell. For example, a description spanning three lines might appear as three distinct rows in the OCR output, with column alignment lost. Extraction models must recognize that these lines belong to a single logical cell and reconstruct the original cell boundaries from visual or positional cues.

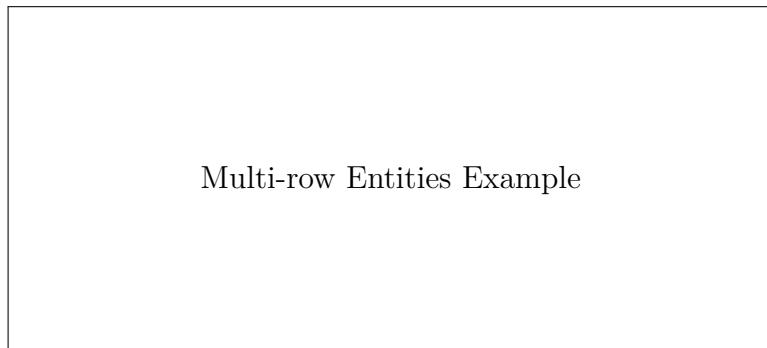


Figure 3: Example of a cell containing multiple lines of text.

Exact duplicates. Production documents sometimes contain intentionally repeated records. In insurance contexts, the same incident may appear multiple times due to amendments, re-openings, or reporting across multiple policy periods. Unlike data entry errors, these duplicates are semantically meaningful and must be preserved in the extracted output. Many

extraction pipelines include deduplication as a post-processing step, which would incorrectly collapse valid duplicate entries. Models must faithfully reproduce the document content without applying implicit deduplication logic.

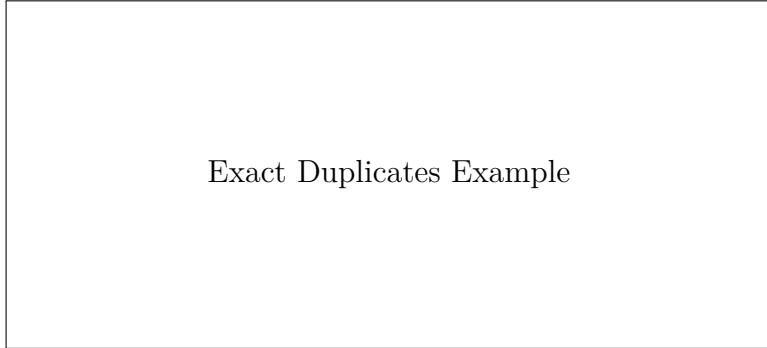


Figure 4: Example of duplicate incident records.

Large documents. Real loss runs and itemized bills routinely contain hundreds of line items. A single trucking company’s annual loss run may list 200–500 incidents; a hospital bill for a complex procedure can exceed 1,000 charge lines. These documents push against context window limits of current LLMs. Even models with 128K+ token windows may exhibit degraded recall on items appearing in the middle of very long documents (the “lost in the middle” phenomenon). We include an extreme tier with 500 incidents per document to measure how extraction quality scales with document length.

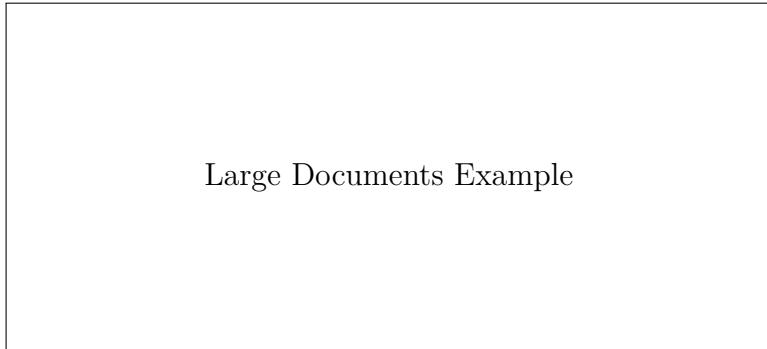


Figure 5: Example of a document with hundreds of incidents.

Multiple tables. Business documents frequently embed auxiliary tables alongside the primary data. A loss run PDF might include a cover page with agent contact information, a summary table of totals by coverage type, or a glossary of status codes—none of which should be extracted as incident records. Models must distinguish the target entity table from these distractors based on schema matching, header recognition, or positional cues. Naive approaches that extract all tabular content will produce false positives from irrelevant tables.

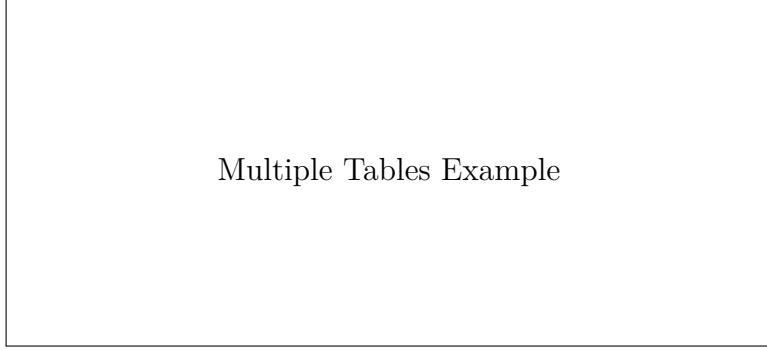


Figure 6: Example of irrelevant tables appearing alongside claims data.

Multi-column layout. Some documents use multi-column layouts to fit more content per page. This creates reading-order ambiguity: should text be read left-to-right across columns or top-to-bottom within each column? Standard OCR often assumes a single-column flow, interleaving content from parallel columns into an incoherent sequence. For example, incident #1’s description might be concatenated with incident #10’s financial data if both appear at the same vertical position in adjacent columns. Correct extraction requires column detection and proper reading-order reconstruction.

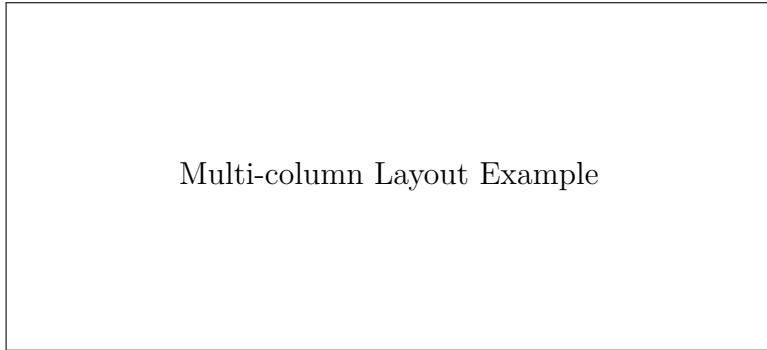


Figure 7: Example of a two-column document layout.

Merged cells. Tables often use merged cells for visual grouping. A common pattern is merging the company name cell across all incidents belonging to that company, or merging a coverage type header across its associated rows. In the rendered PDF, these appear as single cells spanning multiple rows or columns. OCR may report the merged cell’s content only once, leaving subsequent rows with empty values for that field. Extraction models must propagate the merged value to all spanned rows, recognizing that an empty cell indicates inheritance from above rather than a missing value.

3.4 Dataset scale

The released dataset (`longlistbench-v1`, version 1.0.1; see `benchmarks/claims/metadata.json`) contains 80 PDFs (40 detailed, 40 table) with 6,828 incident rows. Difficulty tiers

Merged Cells Example

Figure 8: Example of a table with merged and spanning cells.

are configured as 15 easy instances (10 claims/PDF), 12 medium (25 claims/PDF), 8 hard (50 claims/PDF), and 5 extreme (100 claims/PDF nominal). In the extreme tier, enabling `large_doc` expands each document to 500 incidents. Enabling `duplicates` injects additional duplicate rows (up to 5 per document), causing the number of incident rows to exceed the nominal tier size.

Across the 80 documents, the most common injected issues are multi-row entities (62/80), page breaks (56/80), duplicates (56/80), and multiple tables (40/80).

4 Evaluation

We evaluate systems on the task of extracting a list of incident records from the OCR transcript of each PDF. The benchmark provides both the OCR transcript and a structured JSON ground truth for each document.

4.1 OCR transcription

All PDFs are converted to Markdown using a Gemini vision model (`benchmarks/ocr_claims_pdfs.py`). Each page is rendered to an image and transcribed with a system prompt that emphasizes preserving layout, spacing, and tables. In particular, tables are emitted in a CSV-like form inside Markdown, and the output is concatenated across pages.

4.2 LLM extraction protocol

We provide a lightweight, zero-shot evaluation harness (`benchmarks/evaluate_models.py`) that applies the same extraction prompt to multiple LLM providers and requires the model to return a JSON list of incident objects conforming to the full incident schema. The prompt includes a JSON Schema serialization of the target Pydantic model and is executed at temperature 0. Where supported, we request native structured outputs (e.g., response schemas) to reduce formatting errors.

To ensure schema conformance, model outputs are validated and normalized against a Pydantic schema before scoring. Predictions are stored as `{sample}_{model}_predicted.json`, and aggregate reports are exported as `evaluation_report.json` and `evaluation_report.md`.

4.3 Chunking and merging for long documents

Hard and extreme documents can contain hundreds of incidents, and the OCR transcript may exceed practical context limits. The evaluation harness therefore supports chunked extraction: the OCR text is split into overlapping chunks using simple incident-number markers, targeting at most eight incidents per chunk. Each chunk is extracted independently, and chunk-level predictions are merged by normalized incident identifier, preferring non-empty fields and combining nested financial breakdown subfields.

4.4 Report regeneration and validation

To support reproducible analysis, the harness can regenerate summary reports offline from saved prediction files and optionally reuse extraction-time values from a previous report. A companion checker script (`benchmarks/check_evaluation_report.py`) recomputes metrics from the saved predictions and the golden data, and flags schema violations or report inconsistencies.

4.5 Field-level matching and metrics

For scoring we use the incident number as the record identifier. Incident numbers are normalized by stripping common prefixes (e.g., #, Incident #). Let G be the list of ground-truth records and P be the list of predicted records. We compute micro precision/recall/F1 over field-value pairs, after canonicalizing each incident under the schema.

Canonicalization strips whitespace from strings, maps empty optional strings to null, sorts claimant lists, and rounds monetary values in nested financial breakdowns to two decimal places. Metrics are computed per document and then averaged across documents for tier- and format-level summaries.

For each incident, we flatten its fields into a multiset of canonicalized triples (`incident_id`, `field_path`, `value`) (including nested financial breakdown fields). We then define:

$$\text{found} = |\mathcal{F}(G) \cap \mathcal{F}(P)|, \quad (1)$$

$$\text{recall} = \frac{\text{found}}{|\mathcal{F}(G)|}, \quad (2)$$

$$\text{precision} = \frac{\text{found}}{|\mathcal{F}(P)|}, \quad (3)$$

$$F1 = \frac{2 \text{precision recall}}{\text{precision} + \text{recall}}, \quad (4)$$

where $\mathcal{F}(\cdot)$ denotes the multiset of flattened field-value pairs across incidents. We additionally report missing and extra incident identifiers and count exact record matches for incidents whose canonicalized objects match exactly.

5 Results

We summarize results for (i) OCR fidelity and (ii) baseline extraction performance. OCR coverage and LLM baseline numbers are produced by the released scripts in the repository.

Table 1: OCR identifier coverage on the full dataset (80 documents).

Identifier	Mean coverage	Min coverage
Incident number	100.0%	100.0%
Reference number	100.0%	100.0%

Table 2: Zero-shot LLM baseline results across the full benchmark (80 documents) under schema-conformant, field-level scoring (computed from `benchmarks/results_*_all/evaluation_report.json`).

Model	Samples	Avg Recall	Avg Precision	Avg F1
Gemini 2.5	80	80.4%	83.4%	81.9%
GPT-4o	80	78.3%	82.0%	80.0%
GPT-5.2	80	76.8%	79.6%	78.1%

5.1 OCR identifier coverage

Using `benchmarks/validate_ocr_vs_golden.py` we measure how often key identifiers from the ground truth appear verbatim in the OCR transcript. Across the full dataset (80 OCR transcripts), incident numbers and reference numbers exhibit 100% coverage (mean and minimum). These results indicate that, for primary identifiers, our OCR step rarely drops information and that most downstream failures are attributable to extraction rather than transcription.

5.2 Zero-shot LLM extraction baseline

We evaluate three LLMs using the shared prompt and evaluation harness in `benchmarks/evaluate_models`. We report schema-conformant, field-level scoring across the full benchmark (80 documents: 40 detailed, 40 table) using the released per-tier evaluation reports (`benchmarks/results_*_all/evaluation_report.*`). Averaged across all documents, Gemini 2.5 achieves 81.9% average F1 (80.4% recall, 83.4% precision), GPT-4o achieves 80.0% average F1 (78.3% recall, 82.0% precision), and GPT-5.2 achieves 78.1% average field-level F1 (76.8% recall, 79.6% precision) (Table 2). Across all models, the detailed format is substantially easier than the table format (Table 3), and performance varies meaningfully across difficulty tiers (Table 4).

Qualitatively, errors often manifest as local field-level deviations (e.g., missing optional strings, numeric drift in financial breakdowns, or small identifier formatting mistakes) spread across an otherwise correct long list.

These findings suggest that recovering identifiers is largely deterministic under our OCR pipeline, while the main open challenge for long-list extraction is robustly segmenting and populating full per-incident records under layout disruptions (page breaks, multi-column order, irrelevant tables, merged cells) and scale (hundreds of incidents).

Table 3: Baseline F1 by document format aggregated across all tiers (`benchmarks/results_*_all/evaluation_report.json`).

Model	Detailed F1	Table F1
Gemini 2.5	89.8%	73.9%
GPT-4o	89.3%	70.8%
GPT-5.2	83.5%	72.8%

Table 4: Baseline F1 by difficulty tier (average across documents within each tier; `benchmarks/results_*_all/evaluation_report.json`).

Tier	Samples	Gemini 2.5 F1	GPT-4o F1	GPT-5.2 F1
Easy	30	85.1%	82.2%	80.2%
Medium	24	80.5%	79.7%	76.7%
Hard	16	78.1%	76.6%	76.0%
Extreme	10	81.6%	80.0%	78.9%

6 Limitations and Future Directions

6.1 Limitations

LongListBench is designed to be a practical benchmark for long-list extraction in semi-structured business documents, but it has several limitations.

- **Synthetic generation:** while the schema and failure modes are motivated by production documents, the records and layouts are programmatically generated. This may under-represent rare formatting conventions and highly idiosyncratic carrier templates.
- **OCR stack:** OCR transcripts are produced by a single vision-language model with a fixed prompt. Classical OCR systems and alternative prompts may yield different error distributions.
- **Domain scope:** the current instances are inspired by recurring patterns observed in real-world claims documents. Generalization to other long-list domains (invoices, purchase orders, medical billing, financial statements) should be validated.
- **Metric granularity:** while we report schema-conformant, field-level scoring, the metric does not capture semantic equivalence or provide exact-match analyses, and it remains challenging to evaluate duplicates as first-class entities.

6.2 Future directions

We see multiple immediate extensions that would strengthen LongListBench and increase its utility.

- **Per-field analysis:** report more detailed per-field accuracy for the full incident schema, and explicitly evaluate duplicate detection and normalization.
- **Exact-match analysis:** evaluate the ability of models to produce exact matches for extracted fields.
- **Broader OCR conditions:** add scans with blur/noise, different DPI settings, and non-LLM OCR baselines.
- **Broader document families:** add additional templates and long-list document types beyond the claims-inspired formats in the current release.
- **Scalable extraction protocols:** benchmark chunking, retrieval-augmented extraction, and layout-aware reconstruction strategies for the extreme tier (500 incidents).

7 Conclusion

LongListBench targets a persistent gap in document understanding evaluation: extracting long lists of repeated entities from semi-structured business documents under realistic layout and OCR noise. We presented a benchmark construction pipeline that produces paired (PDF, OCR, JSON) artifacts and systematically injects common long-list failure modes.

7.1 Summary

Our main contributions are:

- A reproducible benchmark generation pipeline for semi-structured documents with long incident lists spanning two formats and four difficulty tiers.
- A taxonomy of seven problem types that frequently break long-list extraction systems, including duplicates, page breaks, multi-row entities, multi-column layout, and merged cells.
- An evaluation harness and baseline results that quantify the gap between near-perfect OCR identifier retention and imperfect end-to-end extraction.

7.2 Practical takeaways

We intend LongListBench to be useful as a measurement tool for both research and engineering workflows. Two practical takeaways are worth emphasizing. First, identifier retention in OCR is near-perfect (Table 1), so most end-to-end failures should be attributed to downstream parsing, segmentation, and field population rather than transcription. Second, even with schema-conformant structured outputs and a shared prompt, field-level extraction across the full benchmark remains materially below perfect (Table 2), with a large gap between detailed and table formats (Table 3) and meaningful variation across difficulty tiers (Table 4), indicating substantial headroom for methods that explicitly model reading order, table structure, and long-range consistency.

7.3 Recommended reporting

For comparability across papers and systems, we recommend that LongListBench results report (i) OCR identifier coverage, (ii) schema-conformant field-level precision/recall/F1 under the released evaluator, and (iii) the extraction protocol used for long documents (e.g., full-context vs chunking, chunk sizes, and merge strategy). The extreme tier, in particular, is intended to stress scaling behavior: methods that succeed on short lists may fail due to context limits, brittle segmentation, or accumulated small errors across hundreds of records.

7.4 Future Work

We view the benchmark as a foundation for studying scalable, layout-robust extraction. Immediate next steps include improved handling of duplicates and merged cells, and evaluation of methods that can reliably extract hundreds of incidents in a single document (Section 6).

Acknowledgments

We thank Ben Perryman and Colin for helpful feedback on early drafts and for reviewing the manuscript.

References

- [1] Y. Xu, M. Li, L. Cui, S. Huang, F. Wei, and M. Zhou, *Layoutlm: Pre-training of text and layout for document image understanding*, 2020. doi: [10.1145/3394486.3403172](https://doi.org/10.1145/3394486.3403172). arXiv: [1912.13318 \[cs.CL\]](https://arxiv.org/abs/1912.13318). [Online]. Available: <https://arxiv.org/abs/1912.13318>.
- [2] G. Kim et al., *Ocr-free document understanding transformer*, 2022. arXiv: [2111.15664 \[cs.LG\]](https://arxiv.org/abs/2111.15664). [Online]. Available: <https://arxiv.org/abs/2111.15664>.
- [3] G. Jaume, H. K. Ekenel, and J.-P. Thiran, *Funsd: A dataset for form understanding in noisy scanned documents*, 2019. arXiv: [1905.13538 \[cs.IR\]](https://arxiv.org/abs/1905.13538). [Online]. Available: <https://arxiv.org/abs/1905.13538>.
- [4] Z. Huang et al., *Icdar2019 competition on scanned receipt ocr and information extraction*, 2021. doi: [10.1109/ICDAR.2019.00244](https://doi.org/10.1109/ICDAR.2019.00244). arXiv: [2103.10213 \[cs.AI\]](https://arxiv.org/abs/2103.10213). [Online]. Available: <https://arxiv.org/abs/2103.10213>.
- [5] Š. Šimsa et al., *Docile benchmark for document information localization and extraction*, 2023. arXiv: [2302.05658 \[cs.CL\]](https://arxiv.org/abs/2302.05658). [Online]. Available: <https://arxiv.org/abs/2302.05658>.
- [6] Z. Wang, Y. Zhou, W. Wei, C.-Y. Lee, and S. Tata, “Vrdu: A benchmark for visually-rich document understanding,” in *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, arXiv:2211.15421, 2023. doi: [10.1145/3580305.3599929](https://doi.org/10.1145/3580305.3599929). [Online]. Available: <https://arxiv.org/abs/2211.15421>.

Appendix A: Evaluation Schemas

A recurring challenge with existing document extraction benchmarks is incomplete or ambiguous schema documentation. Without clear specifications, it can be difficult to understand expected field formats, handling of optional values, or normalization rules. Researchers attempting to reproduce results must often reverse-engineer these details from examples or evaluation scripts, leading to inconsistent implementations and incomparable metrics.

To ensure full reproducibility, we provide complete schemas with explicit type annotations, default values, and field descriptions. All schemas are implemented as Pydantic models, enabling automatic JSON Schema generation and runtime validation. The evaluation script validates model outputs against these schemas before scoring, ensuring that format errors are caught early rather than silently degrading metrics.

A.1 Financial Breakdown Schema

Each incident contains four financial breakdown objects (`bi`, `pd`, `lae`, `ded`) with the following structure:

Listing 1: FinancialBreakdown schema

```
1 class FinancialBreakdown(BaseModel):
2     reserve: float = 0.0
3         # Amount reserved for potential payout
4     paid: float = 0.0
5         # Amount already paid
6     recovered: float = 0.0
7         # Amount recovered (e.g., deductible)
8     total_incurred: float = 0.0
9         # Reserve + Paid - Recovered
```

A.2 Loss Run Incident Schema

The primary entity schema representing a single insurance claim incident:

Listing 2: LossRunIncident schema

```
1 class LossRunIncident(BaseModel):
2     # Identifiers
3     incident_number: str
4         # Incident number (e.g., #12345)
5     reference_number: str
6         # Reference ID (e.g., L240123)
7
8     # Company information
9     company_name: str
10        # Trucking company name
11     division: str = "General"
12        # Company division
```

```

13     insured: str
14         # Insured party name
15     agency: Optional[str] = None
16         # Insurance agency name
17
18     # Policy details
19     policy_number: str
20         # Policy identifier
21     policy_state: str
22         # Policy state abbreviation
23     coverage_type: str
24         # Coverage type (Liability, Physical Damage,
25         # Inland Marine, Cargo)
26     status: str
27         # Open or Closed
28
29     # Incident details
30     description: str
31         # Detailed incident description
32     cause_code: Optional[str] = None
33         # Internal cause code
34     date_of_loss: str
35         # Date incident occurred
36     date_reported: str
37         # Date reported to insurance
38     loss_state: str
39         # State where loss occurred
40
41     # Personnel
42     handler: str = "Claims Adjuster"
43         # Claims handler
44     driver_name: Optional[str] = None
45         # Driver name at time of incident
46     claimants: list[str] = []
47         # List of claimants
48
49     # Vehicle
50     unit_number: Optional[str] = None
51         # Vehicle/truck unit ID
52
53     # Financial breakdowns
54     bi: FinancialBreakdown
55         # Bodily Injury
56     pd: FinancialBreakdown
57         # Property Damage
58     lae: FinancialBreakdown
59         # Loss Adjustment Expense

```

```

60     ded: FinancialBreakdown
61         # Deductible
62
63     # Notes
64     adjuster_notes: Optional[str] = None
65         # Additional adjuster notes

```

A.3 Extraction Output Schema

Models are expected to return a JSON object matching the following structure:

Listing 3: LossRunExtraction schema

```

1 class LossRunExtraction(BaseModel):
2     incidents: list[LossRunIncident]

```

A.4 Field Scoring Rules

During evaluation, fields are normalized and compared as follows:

- **String fields:** Trimmed of whitespace. Optional string fields treat empty strings as `null`.
- **Numeric fields:** Rounded to two decimal places. Negative zero is normalized to zero.
- **List fields:** Sorted alphabetically for comparison.
- **Financial breakdowns:** Each sub-field (`reserve`, `paid`, `recovered`, `total_incurred`) is scored independently.

The evaluation computes field-level precision, recall, and F1 by flattening each incident into (`incident_id`, `field_path`, `value`) tuples and comparing predicted tuples against ground truth.